

[COM4513-6513] Lab 2: Language modelling

Instructor: Nikos Aletras

Teaching Assistants: George Chrysostomou, Hardy and Zeerak Waseem

The goal of this lab is to implement three language models to perform sentence completion, i.e. given a sentence with a missing word to choose the correct one from a list of candidate words. The way to use a language model for this problem is to consider a possible candidate word for the sentence at a time and then ask the language model which version of the sentence is the most probable one.

The sentences to be completed together with the candidate words are in this file: [questions.txt](#). The word to be completed is denoted with ‘_____’ while the pair of candidate words is at the end of the line (e.g. weather/whether). The character ‘.’ between the sentence and the candidates is not part of the sentence. To apply a language model on a sentence for a given candidate word, you just need to replace ‘_____’ with the candidate word.

The texts to train your language models are in this file: [news-corpus-500k.txt](#) (70MB), which is a small subset of the [1 Billion Word Benchmark](#). The text has already been tokenized and split into sentences (each line represents a sentence).

You will implement three language models for this task:

- unigram (2 marks; 1 implementation, 1 report)
- bigram (3 marks; 2 implementation, 1 report)
- bigram with add-1 smoothing, i.e. Laplace (3 marks; 2 implementation, 1 report)

This lab will be marked out of 8 (i.e. 8% of your final grade in the module).

Your code should be executable by running:

```
python3 lab2.py news-corpus-500k.txt questions.txt
```

and present for every question in ‘questions.txt’ the result of each of the three language models. You are advised to lowercase the text and remove punctuation (hint: Python’s `string` module defines punctuation). Please make sure to mention if you’ve used Windows (not recommended) to write and test your code. Make sure your code is Python 3 compatible. See a Python 3 [cheatsheet](#) and [Lab 0](#) if you need help with basic functionality. You are not allowed to use any third-party libraries (e.g. NLTK, Spacy, etc.) to train your language models.

Tip: You can use a dictionary data structure to store the probabilities between unigrams, and unigrams and bigrams, e.g. $p[\text{‘cat’}]$ ($P(\text{cat})$) or $p[\text{‘sat’, ‘cat’}]$ ($P(\text{sat}|\text{cat})$). Remember to append start (`<s>`) and end (`<\s>`) symbols before the first and last words of a sentence respectively.

Your report (2 pages max length) should be submitted as a PDF document containing:

- the accuracy of each model on completing the sentences considering that:
 - in case a language model returns only 0 probabilities, its answer should be considered incorrect
 - in case of a tie with non-zero probabilities, its answer should be considered half-correct
- a discussion of these results, e.g. are they expected?

Optional bonus: Try to implement the language models using: (1) Numpy 2-d arrays (1.5 mark) and (2) Scipy sparse matrices (1.5 mark) instead of Python dictionaries. Are these implementations faster than using Python dictionaries? How big are in the memory using float64 or float32 precision? (Note that these extra marks will be added to your final mark for the lab but that will be capped at 8).

Note that you can use an extra page in your report if you decide to implement the optional bonus.

The deadline for this assignment is the beginning of the Week 5 lab and it needs to be submitted via MOLE. Standard departmental penalties for lateness will be applied.